




Deep Learning Tutorial with PyTorch

Sunglok Choi, Assistant Professor, Ph.D.
Dept. of Computer Science and Engineering, SEOULTECH
sunglok@seoultech.ac.kr | <https://mint-lab.github.io/>


Why Deep Learning?

- **AlphaGo** beat Sedol Lee in a five-game Go match. (2016)

● ALPHAGO
00:08:32



AlphaGo
Google DeepMind




● LEE SEDOL
00:00:27



Why Deep Learning?

- **ChatGPT** broke the [Turing test](#). (2022)
 - Reference) Nature, Vol. 619, 2023 [DOI](#)




Why Deep Learning?

- ChatGPT gives *exact answers* to my *questions*, but **Google** retrieves the *nearest webpages* to my *queries*.



Why Deep Learning?



The image is a screenshot of a Google Scholar search results page for the journal 'Nature'. The browser's address bar shows the URL 'scholar.google.com/citations?hl=en&vq=en&view_op=list_hcore&venue=H--JoiVp8x8J.2020'. The search results are sorted by citation count. The top three results are highlighted with red rectangular boxes. These three papers are all related to deep learning and the game of Go. The other papers listed are about protein-coding genetic variation, human genetic variation, skin cancer classification, and perovskite materials.

Title / Author	Cited by
Deep learning. Y LeCun, Y Bengio, G Hinton Nature 521 (7553), 436	<u>27375</u>
Human-level control through deep reinforcement learning. V Mnih, K Kavukcuoglu, D Silver, AA Rusu, J Veness, MG Bellemare, ... Nature 518 (7540), 529-533	<u>10394</u>
Mastering the game of Go with deep neural networks and tree search. D Silver, A Huang, CJ Maddison, A Guez, L Sifre, G van den Driessche, ... Nature 529 (7587), 484	<u>7698</u>
Analysis of protein-coding genetic variation in 60,706 humans. M Lek, KJ Karczewski, EV Minikel, KE Samocha, E Banks, T Fennell, ... Nature 536 (7616), 285-291	<u>6387</u>
A global reference for human genetic variation. A Auton, LD Brooks, RM Durbin, EP Garrison, HM Kang, JO Korb, ... Nature 526 (7571), 68-74	<u>6011</u>
Dermatologist-level classification of skin cancer with deep neural networks. A Esteva, B Kuprel, RA Novoa, J Ko, SM Swetter, HM Blau, S Thrun Nature 542 (7639), 115-118	<u>4076</u>
Compositional engineering of perovskite materials for high-performance solar cells. NJ Jeon, JH Noh, WS Yang, YC Kim, S Ryu, J Seo, SI Seok Nature 517 (7535), 476-480	<u>4062</u>
Mastering the game of Go without human knowledge. D Silver, J Schrittwieser, K Simonyan, I Antonoglou, A Huang, A Guez, ... Nature 550 (7676), 354	<u>3668</u>

Why Deep Learning?

- Yoshua Bengio, Geoffrey Hinton, and Yann LeCun won the **Turing Award (2018)**.
 - Note) [Chronological listing of A.M. Turing Award Winners](#)



Yoshua Bengio




Geoffrey Hinton



Yann LeCun



Why Deep Learning?

- **AlexNet** won the [ILSVRC \(ImageNet Large Scale Visual Recognition Competition\)](#) 2012 (top-5 error: **16.4%**).
 - An **8-layer** neural network (approx. 61 million parameters) with 2 x GPUs
 - [ReLU \(rectified linear unit\)](#) is used for the vanishing gradient problem and speed-up.
 - [Data augmentation](#) (more data) and [dropout](#) (regularization) was used to solve the overfitting problem.
 - Local response normalization (~ [batch normalization](#)) is used for stable training and generalization.




What is Deep Learning?

- *Deep learning* is machine learning with **deep neural networks** (shortly DNN).
 - e.g. **ResNet** ([residual neural network](#))
 - A **152-layer** network with skip connections
 - The winner of ImageNet Challenge 2015 (top-5 error: **3.57%**)



What is Deep Learning?



- *Deep learning* is **representation learning** (a.k.a. [feature learning](#)).
 - e.g. Task: Draw a line to separate the **red line** and **blue line**.



What is Deep Learning?


- *Deep learning* is **representation learning** (a.k.a. [feature learning](#)).
 - e.g. Task: Draw a line to separate the **red circles** and **green circles**.
 - Note) Try it at [ConvnetJS](#) or [TensorFlow Playground](#).

```
layer_defs = [];  
layer_defs.push({type:'input', out_sx:1, out_sy:1, out_depth:2});  
layer_defs.push({type:'fc', num_neurons:6, activation: 'tanh'});  
layer_defs.push({type:'fc', num_neurons:2, activation: 'tanh'});  
layer_defs.push({type:'softmax', num_classes:2});  
  
net = new convnetjs.Net();  
net.makeLayers(layer_defs);  
  
trainer = new convnetjs.SGDTrainer(net, {learning_rate:0.01, momentum:0.1, batch_size:10, l2_decay:0.001});
```



What is Deep Learning?

- *Deep learning* is **representation learning** (a.k.a. [feature learning](#)).



What is Deep Learning?

- *Deep learning* is **scalable**.




Table of Contents

- **Introduction**
 - Why deep learning?
 - What is deep learning?
- **PyTorch**
 - Why PyTorch?
 - What is PyTorch?
 - Learning PyTorch with Practice
- **Neural Network (NN)**
- **Convolutional Neural Network (CNN)**
- **Recurrent Neural Network (RNN)**
- **Transformer**

Why PyTorch?


- 10 reasons why PyTorch is the DL framework of the future

(by Dhiraj K, September 18th, 2019)

1. PyTorch is Pythonic
2. Easy to learn
3. Higher developer productivity
4. Easy debugging
5. Data parallelism
6. Dynamic computational graph support
7. Hybrid front-end
8. Useful libraries
9. ONNX ([Open Neural Network Exchange](#)) support
10. Cloud support

Why PyTorch?

- Popularity (vs. TensorFlow) in academic communities



What is PyTorch?



- [PyTorch](#) (2016) is an open source deep learning library based on Python.
 - It originated from [Torch](#) (2002) based on the Lua script language.
 - It is primarily developed by Facebook's AI Research lab (FAIR).
- (From my point of view) PyTorch is a **NumPy extension** with supports of
 - + An n-dimensional array with GPU acceleration and automatic differentiation
 - + Useful modules for neural networks
- [My useful lists on deep learning and PyTorch](#)

Practice) PyTorch Installation

- Please follow [PyTorch's installation instruction](#) for your system.
 - Note) If you want GPU acceleration, please install the matched version of CUDA in advance. Please visit [CUDA Toolkit Archive](#) to download a specific version of CUDA.

PyTorch Build	Stable (1.8.1)		Preview (Nightly)	
Your OS	Linux	Mac	Windows	
Package	Conda	Pip	LibTorch	Source
Language	Python		C++ / Java	
Compute Platform	CUDA 10.2	CUDA 11.1	ROCm 4.0 (beta)	CPU
Run this Command:	NOTE: 'conda-forge' channel is required for cudatoolkit 11.1 <code>conda install pytorch torchvision torchaudio cudatoolkit=11.1 -c pytorch -c conda-forge</code>			

- Note) You can use [Google Colab](#) without local installation of PyTorch.

Note) Python Virtual Environment

- A [Python virtual environment](#) is an isolated Python runtime environment that contains a particular version of Python and its additional packages.
 - Note) [Virtual machine](#): A computer entirely implemented by software
- Why virtual environments?
 - We can run Python applications with different versions of packages (e.g. PyTorch 0.9 and PyTorch 1.8) together.
- Tool #1) [venv](#) in the Python Standard Library
- Tool #2) [conda](#) in Anaconda
 - *Conda* is an open source package/dependency/environment management system for programming languages such as Python, R, Java, JavaScript, C/C++, and more.

– Usage

<code>conda create --name venv_name python=3.6</code>	Create a virtual environment
<code>conda create --name venv_dst --clone venv_src</code>	Copy a virtual environment
<code>conda activate venv_name</code>	Activate the virtual environment
<code>conda deactivate</code>	Deactivate the virtual environment
<code>conda env remove --name venv_name</code>	Remove the virtual environment
<code>conda env list</code>	List all virtual environments
<code>conda list</code>	List all installed packages

Learning PyTorch with Practice

(From my point of view) PyTorch is a **NumPy extension** with supports of

- + An n-dimensional array with GPU acceleration and automatic differentiation

so-called a **tensor** (Note: A matrix is a second-order tensor.)

- + Useful modules for neural networks

- Practice with tensors

- Creating a tensor
- Reshaping a tensor
- Line fitting from two points
- CPU vs. GPU-acceleration
- Automatic differentiation (so called *Autograd*)

- Practice with useful modules for neural networks

- Gradient descent by hands and `torch.optim`
- More examples with DNNs, CNNs, and RNNs.

Practice) Creating a Tensor (1/2)

```
import numpy as np
import torch
```

1. Create a tensor from a composite data

```
x = np.array([[3, 29, 82], [10, 18, 84]])
```

```
y = torch.tensor(x)
```

```
print(y.ndim, y.dim())
```

```
# 2
```

Note) x.ndim

```
print(y.nelement())
```

```
# 6
```

Note) x.size

```
print(y.shape, y.size())
```

```
# torch.Size([2, 3])
```

Note) x.shape

```
print(y.dtype)
```

```
# torch.int32
```

Note) x.dtype

2. Create a tensor using initializers

```
p = torch.rand(3, 2)
```

```
# Try zeros, ones, eyes, empty, arange, linspace,
```

```
q = torch.zeros_like(p)
```

```
# and their ..._like
```

```
print(p.dtype)
```

```
# torch.float32
```

```
print(q.shape)
```

```
# torch.Size([3, 2])
```

3. Interpret as a tensor (generating only a view)

```
z = torch.as_tensor(x)
```

```
# Or torch.from_numpy(x) Note) np.asarray()
```

```
x[-1,-1] = 86
```

```
print(z[-1])
```

```
# tensor([10, 18, 86], dtype=torch.int32)
```

Practice) Creating a Tensor (2/2)

4. Access elements

```
print(y[:,1])          # tensor([29, 18])
print(y[0,0])          # tensor(3)
print(y[0,0].item())   # 3
```

Note) `x[0,0] == 3`

5. CUDA tensors

```
if torch.cuda.is_available():
    print(y.device)      # 'cpu'
    y_cuda = y.cuda()   # Or y.to('cuda')
    print(y_cuda.device) # 'cuda:0'
    y_cpu = y_cuda.cpu() # Or y_cuda.to('cpu')
    print(y_cpu.device)  # 'cpu'

    x_cpu = y_cpu.numpy() # Or np.array(y_cpu)
    x_cuda = y_cuda.numpy() # Error!
```

Practice) Reshaping a Tensor (1/2)

```
import numpy as np
import torch
```

```
x = np.array([[[29, 3], [18, 10]], [[27, 10], [12, 5]]])
```

```
y = torch.tensor(x)
```

```
print(y.ndim)           # 3
```

```
print(y.shape)          # torch.Size([2, 2, 2])
```

```
p = y.view(-1)          # tensor([29, 3, 18, 10, 27, 10, 12, 5])
```

```
print(p.shape, p)       # torch.Size([8])
```

```
q = y.view(1, -1)       # tensor([[29, 3, 18, 10, 27, 10, 12, 5]])
```

```
print(q.shape, q)       # torch.Size([1, 8])
```

```
r = y.view(2, -1)       # tensor([[29, 3, 18, 10], [27, 10, 12, 5]])
```

```
print(r.shape, r)       # torch.Size([2, 4])
```

Of course, 'reshape' is also supported and the same as 'view'.

```
s = y.reshape(2, -1, 1) # tensor([[[29], [3], [18], [10]], [[27], [10], [12], [5]]])
```

```
print(s.shape, s)       # torch.Size([2, 4, 1])
```

```
ss = s.squeeze(2)       # Note) s.squeeze(0) and s.squeeze(1) have no effect.
```

```
print(ss)               # tensor([[29, 3, 18, 10], [27, 10, 12, 5]])
```

```
print(ss.shape)         # torch.Size([2, 4])
```

```
u0 = ss.unsqueeze(0)    # tensor([[[29, 3, 18, 10], [27, 10, 12, 5]]])
```

```
print(u0.shape, u0)     # torch.Size([1, 2, 4])
```

```
u1 = ss.unsqueeze(1)    # tensor([[ [29, 3, 18, 10] ], [ [27, 10, 12, 5] ]])
```

```
print(u1.shape, u1)     # torch.Size([2, 1, 4])
```

```
u2 = ss.unsqueeze(2)    # tensor([[[29], [3], [18], [10]], [[27], [10], [12], [5]]])
```

```
print(u2.shape, u2)     # torch.Size([2, 4, 1])
```

29 [0,0,0]	3 [0,0,1]
18 [0,1,0]	10 [0,1,1]

27 [1,0,0]	10 [1,0,1]
12 [1,1,0]	5 [1,1,1]

Practice) Reshaping a Tensor (2/2)

```
# Switch indices each other, (i, j) to (j, i)
t_021 = y.transpose(1, 2)          # tensor([[[29, 18],
print(t_021, t_021.is_contiguous()) #          [ 3, 10]], ... ]) False
c_021 = t_021.contiguous()         # tensor([[[29, 18],
print(c_021, c_021.is_contiguous()) #          [ 3, 10]], ... ]) True
t_102 = y.transpose(0, 1)          # tensor([[[29,  3],
print(t_102, t_102.is_contiguous()) #          [27, 10]], ... ]) False

# Assign indices
t_201 = y.permute(2, 0, 1)          # tensor([[[29, 18],
print(t_201, t_201.is_contiguous()) #          [27, 12]], ... ]) False

# Note) Reshaping does not copy contents.
y[0,0,0] = 27
print(p)                            # tensor([27, 3, 18, 10, 27, 10, 12, 5])
print(t_021)                        # tensor([[[27, 18], ...], ...])
print(c_021)                        # tensor([[[29, 18], ...], ...])

# Copy a tensor and detach it from its connected computational graph
z = y.clone().detach()              # Note) x.clone()
y[0,0,0] = 1
print(y)                            # tensor([[[ 1, 3], ...], ...])
print(z)                            # tensor([[[27, 3], ...], ...])
```

29 [0,0,0]	3 [0,0,1]
18 [0,1,0]	10 [0,1,1]

27 [1,0,0]	10 [1,0,1]
12 [1,1,0]	5 [1,1,1]

t(0,1)

29 [0,0,0]	3 [0,0,1]
18 [1,0,0]	10 [1,0,1]

27 [0,1,0]	10 [0,1,1]
12 [1,1,0]	5 [1,1,1]

relocate

29 [0,0,0]	3 [0,0,1]
27 [0,1,0]	10 [0,1,1]

18 [1,0,0]	10 [1,0,1]
12 [1,1,0]	5 [1,1,1]

Practice) Line Fitting from Two Points

- Find a line which passes through two points, (1,4) and (4,2)
 - The tensor class, `torch.Tensor`, includes member functions of tensor operation and manipulation, and linear algebra (in contrast to `numpy.array`).
 - Note) [PyTorch APIs for torch.Tensor](#)
 - Please remember that a function with ending with an underscore(_) means an in-place function.

```
import torch

A = torch.tensor([[1., 1.], [4., 1.]])
b = torch.tensor([[4.], [2.]])
A_inv = A.inverse()      # Note) np.linalg.inv(A)
print(A_inv.mm(b))       # Note) np.matmul(A_inv, b)
```

- Note) Line fitting with NumPy

```
import numpy as np

A = np.array([[1., 1.], [4., 1.]])
b = np.array([[4.], [2.]])
A_inv = np.linalg.inv(A)
print(np.matmul(A_inv, b)) # [[-0.66666667], [ 4.66666667]]
```


Practice) CPU vs. GPU-acceleration

- Computing time on my laptop
 - cpu : **1.4** [sec] @ Intel i7-7700HQ 2.8GHz
 - cuda: **0.3** [sec] @ NVIDIA GTX 1060

```
import torch
import time

dev_name = 'cuda' if torch.cuda.is_available() else 'cpu' # Try 'cpu'
n = 5000

A = torch.rand(n, n, device=dev_name)
B = torch.rand(n, n, device=dev_name)
start = time.time()
C = A.inverse() * B
elapsed = time.time() - start
print(f'Computing time by {dev_name}: {elapsed:.3f} [sec]')
```



Practice) Automatic Differentiation

- A derivative value, `torch.Tensor.grad`, is available after setting `requires_grad=True` and its forward calculation and backward propagation (a.k.a. [backpropagation](#)).

```
import torch

x = torch.tensor([2.], requires_grad=True)
y = 0.1*x**3 - 0.8*x**2 - 1.5*x + 5.4
y.backward()
print(x.grad) # Derivative: tensor([-3.5000])
```

- Note) Symbolic differentiation with SymPy

```
import sympy as sp

x, y = sp.symbols('x y')
y = 0.1*x**3 - 0.8*x**2 - 1.5*x + 5.4
yd = sp.diff(y, x)
print(yd) # 0.3*x**2 - 1.6*x - 1.5
print(float(yd.subs({x: 2}))) # -3.5
```

Practice) Automatic Differentiation – More Analysis

- Given) $y(x) = x^3$, $z(y) = \log y$
- Q) A derivative value $\frac{\partial z}{\partial x}$ at $x = 5$?

```
import torch

def get_tensor_info(tensor):
    info = []
    for name in ['requires_grad', 'is_leaf', 'retains_grad', 'grad']:
        info.append(f'{name}({getattr(tensor, name, None)})')
    info.append(f'tensor({str(tensor)})')
    return ' '.join(info)

x = torch.tensor(5., requires_grad=True)
y = x ** 3
z = torch.log(y)
print("### Before 'z.backward()')")
print('* x:', get_tensor_info(x))
print('* y:', get_tensor_info(y))
print('* z:', get_tensor_info(z))

y.retain_grad()
z.retain_grad()
z.backward()
print("### After 'z.backward()')")
print('* x:', get_tensor_info(x))
print('* y:', get_tensor_info(y))
print('* z:', get_tensor_info(z))
```

Practice) Automatic Differentiation – More Analysis


- Given) $y(x) = x^3$, $z(y) = \log y$
- Q) A derivative value $\frac{\partial z}{\partial x}$ at $x = 5$?

Before 'z.backward()'

```
* x: requires_grad(True) is_leaf(True) retains_grad(None) grad(None) tensor(tensor(5., requires_grad=True))
* y: requires_grad(True) is_leaf(False) retains_grad(None) grad(None) tensor(tensor(125., grad_fn=<PowBackward0>))
* z: requires_grad(True) is_leaf(False) retains_grad(None) grad(None) tensor(tensor(4.83, grad_fn=<LogBackward>))
```

After 'z.backward()'

```
* x: requires_grad(True) is_leaf(True) retains_grad(None) grad(0.6) tensor(tensor(5., requires_grad=True))
* y: requires_grad(True) is_leaf(False) retains_grad(True) grad(0.008) tensor(tensor(125., grad_fn=<PowBackward0>))
* z: requires_grad(True) is_leaf(False) retains_grad(True) grad(1.) tensor(tensor(4.83, grad_fn=<LogBackward>))
```



Practice) Gradient Descent by Hands

```
f = lambda x: 0.1*x**3 - 0.8*x**2 - 1.5*x + 5.4
viz_range = torch.FloatTensor([-6, 12])
learn_rate = 0.1
max_iter = 100
min_tol = 1e-6
x_init = 12.


# Prepare visualization
xs = torch.linspace(*viz_range, 100)
plt.plot(xs, f(xs), 'r-', label='f(x)', linewidth=2)
plt.plot(x_init, f(x_init), 'b.', label='Each step', markersize=12)
plt.axis((*viz_range, *f(viz_range)))
plt.legend()

x = torch.tensor(x_init, requires_grad=True)
for i in range(max_iter):
    # Derive gradient with Autograd
    if x.grad != None:
        x.grad.zero_()      # Reset gradient tracking
    y = f(x)                # Calculate the function (forward)
    y.backward()            # Calculate the gradient (backward)

    # Run the gradient descent
    xp = x.clone().detach() # Note) xp = x
    with torch.no_grad():   # Disable gradient tracking
        x -= learn_rate*x.grad # Note) x = x - learn_rate*fd(x) is an original code.
                                # x = x - learn_rate*x.grad() does not work!

    # Update visualization for each iteration
    print(f'Iter: {i}, x = {xp:.3f} to {x:.3f}, f(x) = {f(xp):.3f} to {f(x):.3f} (f\'(x) = {x.grad:.3f})')
    lcolor = torch.rand(3).tolist()
    approx = x.grad*(xs-xp) + f(xp)
    plt.plot(xs, approx, '-', linewidth=1, color=lcolor, alpha=0.5)
    xc = x.clone().detach() # Copy 'x' for plotting
    plt.plot(xc, f(xc), '.', color=lcolor, markersize=12)

    # Check the terminal condition
    if abs(x - xp) < min_tol:
        break;
plt.show()
```



Practice) Gradient Descent by torch.optim

```
f = lambda x: 0.1*x**3 - 0.8*x**2 - 1.5*x + 5.4
viz_range = torch.FloatTensor([-6, 12])
learn_rate = 0.1
max_iter = 100
min_tol = 1e-6
x_init = 12.

# Prepare visualization
xs = torch.linspace(*viz_range, 100)
plt.plot(xs, f(xs), 'r-', label='f(x)', linewidth=2)
plt.plot(x_init, f(x_init), 'b.', label='Each step', markersize=12)
plt.axis((*viz_range, *f(viz_range)))
plt.legend()

x = torch.tensor(x_init, requires_grad=True)
optimizer = torch.optim.SGD([x], lr=learn_rate)
for i in range(max_iter):
    # Run the gradient descent with the optimizer
    optimizer.zero_grad()      # Reset gradient tracking
    y = f(x)                   # Calculate the function (forward)
    y.backward()                # Calculate the gradient (backward)
    xp = x.clone().detach()     # Note) xp = x
    optimizer.step()            # Update 'x'

    # Update visualization for each iteration
    print(f'Iter: {i}, x = {xp:.3f} to {x:.3f}, f(x) = {f(xp):.3f} to {f(x):.3f} (f\'(x) = {x.grad:.3f})')
    lcolor = torch.rand(3).tolist()
    approx = x.grad*(xs-xp) + f(xp)
    plt.plot(xs, approx, '-', linewidth=1, color=lcolor, alpha=0.5)
    xc = x.clone().detach()    # Copy 'x' for plotting
    plt.plot(xc, f(xc), '.', color=lcolor, markersize=12)

    # Check the terminal condition
    if abs(x - xp) < min_tol:
        break;
plt.show()
```





Table of Contents


- Introduction
- PyTorch
- **Neural Network (NN)**
 - Perceptron
 - Multi-layer perceptron
 - Activation function
 - Backpropagation
 - Vanishing gradient problem
 - Loss function
 - Example) Iris flower classification
- **Convolutional Neural Network (CNN)**
- **Recurrent Neural Network (RNN)**
- **Transformer**

Neural Network

- An *artificial neural network* (shortly *neural network*, *NN*) is a collection of perceptrons (a.k.a. artificial neurons) and their connection with weights.
 - Inspired by the biological neural networks.
- **Neuron vs. Perceptron**



vs.




$$y = \phi \left(\sum_{i=1}^n w_i x_i + b \right)$$

where ϕ is an activation function (for nonlinearity)
and b is bias.


Neural Network

- **Multi-layer perceptrons (MLP)**
 - Note) Fully-connected (shortly FC) layer



2-layer NN

(1 x hidden layer, 1 x output)







3-layer NN

(2 x hidden layer, 2 x output)

Activation Function

- Activation function imposes **nonlinearity** to a neural network.



Practice) Visualizing Activation Functions

- Visualize activation functions and their derivative functions

- Note) [PyTorch APIs for activation functions](#)

```
import torch
import torch.nn as nn
import matplotlib.pyplot as plt

activation_funcs = [
    {'name': 'Sigmoid',      'func': nn.Sigmoid()},
    {'name': 'Tanh',         'func': nn.Tanh()},
    {'name': 'ReLU',         'func': nn.ReLU()},
    {'name': 'Leaky ReLU',   'func': nn.LeakyReLU(0.1)},
    {'name': 'ELU',          'func': nn.ELU()},
    # Try more activation functions
]

for act in activation_funcs:
    x = torch.linspace(-10, 10, 200, requires_grad=True)
    y = act['func'](x)
    y.sum().backward()

    plt.title(act['name'])
    x_np, y_np, grad = x.detach().numpy(), y.detach().numpy(), x.grad.numpy()
    plt.plot(x_np, y_np, label='$\phi(x)$')
    plt.plot(x_np, grad, label='$\partial \phi(x) / \partial x$')
    plt.grid()
    plt.legend()
    plt.show()
```


Backpropagation

- **Training a neural network (~ optimization)**

Finding weight variables which minimize a cost function as

$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} \frac{1}{N} \sum_{d=1}^N l(\mathbf{y}^d, \hat{\mathbf{y}}^d)$$

where l is a loss function, N is the number of data, and $\hat{\mathbf{y}}^d$ is the d -th target value.




Evaluation with mean squared error

$$MSE = \frac{1}{N} \sum_{d=1}^N (\hat{\mathbf{y}}^d - \mathbf{y}^d)^2$$

where $l(\mathbf{y}^d, \hat{\mathbf{y}}^d) = (\hat{\mathbf{y}}^d - \mathbf{y}^d)^2$

Backpropagation

- Backpropagation is an algorithm for training a neural network.
 - It tries to find the optimal **weight variables** of a NN by **gradient descent**
- **Vanishing gradient problem**
 - During backpropagation, gradient values of a deep NN become close to 0.




Practice) Observing **Vanishing** Gradients

- Gradients of a single-node multi-layer NN


- 1-layer forward: $y = \phi(wx + b) = \phi(x)$ **if $w = 1$ and $b = 0$**

- 1-layer backward: $\frac{\partial y}{\partial x} = \phi'(x)$




- 2-layer forward: $y = \phi(\phi(x))$

- 2-layer backward: $\frac{\partial y}{\partial x} = \phi'(\phi(x))\phi'(x)$ **$\therefore$ chain rule**



- 3-layer forward: $y = \phi(\phi(\phi(x)))$

- 3-layer backward: $\frac{\partial y}{\partial x} = \phi'(\phi(\phi(x)))\phi'(\phi(x))\phi'(x)$



- ...

Practice) Observing **Vanishing** Gradients

- Gradients of a single-node multi-layer NN with the **sigmoid** function


– Note) $\phi'(x) = \phi(x)(1 - \phi(x))$

```
import numpy as np
import matplotlib.pyplot as plt
```

```
f = lambda x: 1 / (1 + np.exp(-x))
df = lambda x: f(x) * (1 - f(x))
```

```
x = np.linspace(-10, 10, 1000)
```

```
plt.plot(x, df(x), label='1-layer')
plt.plot(x, df(f(x))*df(x), label='2-layer')
plt.plot(x, df(f(f(x)))*df(f(x))*df(x), label='3-layer')
plt.plot(x, df(f(f(f(x))))*df(f(f(x)))*df(f(x))*df(x), label='4-layer')
plt.axis((-10, 10, 0, 0.3))
plt.grid()
plt.legend()
plt.show()
```



Practice) Observing Vanishing Gradients


- Gradients of a single-node multi-layer NN with the **ReLU** function

```
import numpy as np
import matplotlib.pyplot as plt
```

```
f = lambda x: x * (x >= 0)
df = lambda x: x >= 0
```

```
x = np.linspace(-10, 10, 1000)
```



```
plt.plot(x, df(x), label='1-layer')
plt.plot(x, df(f(x))*df(x), label='2-layer')
plt.plot(x, df(f(f(x)))*df(f(x))*df(x), label='3-layer')
plt.plot(x, df(f(f(f(x))))*df(f(f(x)))*df(f(x))*df(x), label='4-layer')
plt.axis((-10, 10, 0, 0.3))
plt.grid()
plt.legend()
plt.show()
```



Activation Function (Revisited)

- Why [ReLU](#)?

- It can mitigate the vanishing gradient problem.
- In addition, there is a biological analogue such as neural [action potential](#).



Loss Function

- **Training a neural network (~ optimization)**

Finding weight variables which minimize a cost function as

$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} \frac{1}{N} \sum_{d=1}^N l(\mathbf{y}^d, \hat{\mathbf{y}}^d)$$

where l is a loss function, N is the number of data, and $\hat{\mathbf{y}}^d$ is the d -th target value.

- **Loss function**

- A *loss function* quantifies the gap between the ground truth and prediction.
- e.g. **Mean squared error** (usually for *regression*)

$$l(\mathbf{y}, \hat{\mathbf{y}}) = (\hat{\mathbf{y}} - \mathbf{y})^2$$

where $\hat{\mathbf{y}}$ is the ground truth, $\mathbf{y}$ is prediction, and N is the number of data.

- e.g. **Binary cross entropy error** (usually for *binary classification*)

$$l(\mathbf{y}, \hat{\mathbf{y}}) = -\hat{\mathbf{y}} \log \mathbf{y} - (1 - \hat{\mathbf{y}}) \log(1 - \mathbf{y})$$

In general, $-\sum \hat{\mathbf{y}}_i \log \mathbf{y}_i$ for *multi-class classification*


Loss Function

- **Loss function**

- e.g. **Cross entropy error** (usually for *multi-class classification*)

$$l(\mathbf{y}, \hat{\mathbf{y}}) = - \sum_{i=1}^c \hat{y}_i \log y_i$$

where $\hat{y}_i$ is the one-hot-encoded truth, y_i is the predicted (softmax) confidence



- Note) [PyTorch APIs for loss functions](#)


Practice) Iris Flower Classification (1/3)

```
import torch
import torch.nn as nn
import numpy as np
import matplotlib.pyplot as plt
from sklearn import (datasets, metrics)
from matplotlib.colors import ListedColormap
import time

# 1.1. Load a dataset partially
iris = datasets.load_iris()
iris.data = iris.data[:,0:2]
iris.feature_names = iris.feature_names[0:2]
iris.color = np.array([(1, 0, 0), (0, 1, 0), (0, 0, 1)])

# 1.2. Load the dataset as tensors
dev_name = 'cuda' if torch.cuda.is_available() else 'cpu' # Try 'cpu'
x = torch.tensor(iris.data, device=dev_name).float()
y = torch.tensor(iris.target, device=dev_name).long()

# 2. Define a model
# - Try the different number of hidden layers
# - Try less or more layers with different transfer functions
input_size, output_size = len(iris.feature_names), len(iris.target_names)
model = nn.Sequential(
    nn.Linear(input_size, 4),
    nn.ReLU(),
    nn.Linear(4, output_size),
).to(dev_name)
```



Practice) Iris Flower Classification (2/3)



3. Train the model

```
loss_func = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=0.01) # Try other optimizers
epoch_max = 10000
loss_list = []
start = time.time()
for i in range(epoch_max):
    # Train one iteration
    optimizer.zero_grad()
    output = model(x)
    loss = loss_func(output, y)
    loss.backward()
    optimizer.step()

    # Record the loss
    loss_list.append(loss / len(x))
elapsed = time.time() - start
```

4.1. Visualize the training loss curve

```
plt.title(f'Training Loss (time: {elapsed/60:.2f} [min] @ {dev_name})')
plt.plot(range(1, epochs + 1), loss_list)
plt.xlabel('Epochs')
plt.ylabel('Loss values')
plt.xlim((1, epochs))
plt.grid()
plt.show()
```




Practice) Iris Flower Classification (3/3)

4.2. Visualize training results (decision boundaries)

```
x_min, x_max = iris.data[:, 0].min() - 1, iris.data[:, 0].max() + 1
y_min, y_max = iris.data[:, 1].min() - 1, iris.data[:, 1].max() + 1
xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01), np.arange(y_min, y_max, 0.01))
xy = np.vstack((xx.flatten(), yy.flatten())).T
xy_tensor = torch.from_numpy(xy).float().to(dev_name)
zz = torch.argmax(model(xy_tensor), dim=1).cpu().detach().numpy()
plt.contourf(xx, yy, zz.reshape(xx.shape), cmap=ListedColormap(iris.color), alpha=0.2)
```

4.3. Visualize data with their classification

```
predict = torch.argmax(model(x), dim=1).cpu().detach().numpy()
accuracy = metrics.balanced_accuracy_score(iris.target, predict)
plt.title(f'Fully-connected NN (accuracy: {accuracy:.3f})')
plt.scatter(iris.data[:,0], iris.data[:,1], c=iris.color[iris.target], edgecolors=iris.color[predict])
plt.xlabel(iris.feature_names[0])
plt.ylabel(iris.feature_names[1])
plt.show()
```



Practice) Iris Flower Classification – My Style (1/5)


```
import torch
import torch.nn as nn
import torch.nn.functional as F
import numpy as np
import matplotlib.pyplot as plt
from sklearn import datasets
from matplotlib.colors import ListedColormap
import time

# Define hyperparameters
EPOCH_MAX = 10000
EPOCH_LOG = 1000
OPTIMIZER_PARAM = {'lr': 0.01}
DATA_LOADER_PARAM = { 'batch_size': 50, 'shuffle': True }
USE_CUDA = torch.cuda.is_available()
RANDOM_SEED = 777

# A two-layer NN model
class MyDNN(nn.Module):
    def __init__(self, input_size=2, output_size=3):
        super(MyDNN, self).__init__()
        self.fc1 = nn.Linear(input_size, 4)
        self.fc2 = nn.Linear(4, output_size)

        nn.init.xavier_uniform_(self.fc1.weight)
        nn.init.xavier_uniform_(self.fc2.weight)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```




Practice) Iris Flower Classification – My Style (2/5)

Train a model with the given batches

```
def train(model, batch_data, loss_func, optimizer):  
    model.train() # Notify layers (e.g. Dropout, BatchNorm) that it's now training  
    train_loss, n_data = 0, 0  
    dev = next(model.parameters()).device  
    for batch_idx, (x, y) in enumerate(batch_data):  
        x, y = x.to(dev), y.to(dev)  
        optimizer.zero_grad()  
        output = model(x)  
        loss = loss_func(output, y)  
        loss.backward()  
        optimizer.step()  
  
        train_loss += loss.item()  
        n_data += len(y)  
    return train_loss / n_data
```

Evaluate a model with the given batches

```
def evaluate(model, batch_data, loss_func):  
    model.eval() # Notify layers (e.g. Dropout, BatchNorm) that it's now testing  
    test_loss, n_correct, n_data = 0, 0, 0  
    with torch.no_grad():  
        dev = next(model.parameters()).device  
        for x, y in batch_data:  
            x, y = x.to(dev), y.to(dev)  
            output = model(x)  
            loss = loss_func(output, y)  
            y_pred = torch.argmax(output, dim=1)  
  
            test_loss += loss.item()  
            n_correct += (y == y_pred).sum().item()  
            n_data += len(y)  
    return test_loss / n_data, n_correct / n_data
```



Practice) Iris Flower Classification – My Style (3/5)

```
if __name__ == '__main__':
    # 0. Preparation
    torch.manual_seed(RANDOM_SEED)
    if USE_CUDA:
        torch.cuda.manual_seed_all(RANDOM_SEED)
    dev = torch.device('cuda' if USE_CUDA else 'cpu')

    # 1.1. Load the Iris dataset partially
    iris = datasets.load_iris()
    iris.data = iris.data[:,0:2]
    iris.feature_names = iris.feature_names[0:2]
    iris.color = np.array([(1, 0, 0), (0, 1, 0), (0, 0, 1)])

    # 1.2. Wrap the dataset with torch.utils.data.DataLoader
    x = torch.tensor(iris.data, dtype=torch.float32, device=dev)
    y = torch.tensor(iris.target, dtype=torch.long, device=dev)
    data_train = torch.utils.data.TensorDataset(x, y)
    loader_train = torch.utils.data.DataLoader(data_train, **DATA_LOADER_PARAM)
```

Practice) Iris Flower Classification – My Style (4/5)

```
# 2. Instantiate a model, loss function, and optimizer
model = MyDNN().to(dev)
loss_func = F.cross_entropy
optimizer = torch.optim.SGD(model.parameters(), **OPTIMIZER_PARAM)

# 3. Train the model
loss_list = []
start = time.time()
for epoch in range(1, EPOCH_MAX + 1):
    train_loss = train(model, loader_train, loss_func, optimizer)
    valid_loss, valid_accuracy = evaluate(model, loader_train, loss_func)

    loss_list.append([epoch, train_loss, valid_loss, valid_accuracy])
    if epoch % EPOCH_LOG == 0:
        elapse = time.time() - start
        print(f'{epoch:>6} ({elapse:>6.2f} sec), TrLoss={train_loss:.6f}, VaLoss={valid_loss:.6f}, VaAcc={valid_accuracy:.3f}')
elapse = time.time() - start

# 4.1. Visualize the loss curves
# 4.2. Visualize training results (decision boundaries)
# 4.3. Visualize data with their classification
```

Practice) Iris Flower Classification – My Style (5/5)

```
1000 ( 7.71 sec), TrLoss=0.009773, VaLoss=0.009769, VaAcc=0.707
2000 (15.17 sec), TrLoss=0.009300, VaLoss=0.009281, VaAcc=0.747
3000 (22.85 sec), TrLoss=0.009022, VaLoss=0.009007, VaAcc=0.780
4000 (30.41 sec), TrLoss=0.008800, VaLoss=0.008786, VaAcc=0.800
5000 (37.85 sec), TrLoss=0.008606, VaLoss=0.008594, VaAcc=0.813
6000 (45.49 sec), TrLoss=0.008433, VaLoss=0.008420, VaAcc=0.813
7000 (53.06 sec), TrLoss=0.008274, VaLoss=0.008262, VaAcc=0.813
8000 (60.31 sec), TrLoss=0.008138, VaLoss=0.008118, VaAcc=0.807
9000 (67.91 sec), TrLoss=0.008008, VaLoss=0.007990, VaAcc=0.800
10000 (75.14 sec), TrLoss=0.007928, VaLoss=0.007880, VaAcc=0.800
```

Discussion) CPU vs. GPU

Data as raw tensors vs. Data as DataLoader





Table of Contents


- Introduction
- PyTorch
- Neural Network (NN)
- **Convolutional Neural Network (CNN)**
 - Convolutional layer
 - Pooling layer
 - Dropout
 - Skip connection
 - Example) Digit classification with the MNIST dataset
- **Recurrent Neural Network (RNN)**
- **Transformer**

Convolution

- Discrete convolution


$$(f * g)(k) = \sum_{i=-\infty}^{\infty} f(i)g(k-i)$$

- 1D discrete convolution (e.g. voice)



Convolution

- 1D discrete convolution (continued)




- The size of output vectors

$$\text{output_size} = \left\lfloor \frac{\text{input_size} + 2 \times \text{padding} - (\text{dilation} \times (\text{kernel_size} - 1) + 1)}{\text{stride}} + 1 \right\rfloor$$

Convolution

- 2D discrete convolution (e.g. image)





- A `kernel` and its `output` are also called a `filter` and `feature map` (or `activation map`), respectively.
 - CNN visualization examples: [ConvNetJS](#), [CNN-Explainer](#)
 - e.g. When a kernel is a horizontal Sobel filter,



Convolutional Layer

- A **convolutional layer** (shortly *conv layer*) is a NN layer which uses *convolution* as its feedforward propagation and *kernel* as weight variables.
- (In contrast to a FC layer) A convolutional layer has two important points, **weight sharing** and **local connectivity**.
 - e.g. `input_size = 6`, `output_size = 4`, `kernel_size = 3`
FC layer has 24 weight variables, but conv layer has only 3 weight variables.




Why it works?




Convolutional Layer

- A convolutional layer can have multiple kernels so that its output will be multiple *channels*.
 - e.g. `in_channel_size = 1`, `out_channel_size = 2`, `kernel_size = 3`
 - e.g. `in_channel_size = m`, `out_channel_size = n`, `kernel_size = 3`



of weights (w/o bias): 3×2




of weights (w/o bias): $3 \times 3 \times m \times n$

- Note) The number of weights is independent from input and output size in a convolutional layer.
 - Note) The number of weight (w/o bias) in a FC layer = input size x output size
 - e.g. The left example: $6 \times 8 = 48$


Pooling Layer

- A **pooling layer** performs a non-linear down-sampling operation.
- Especially max pooling is common.
 - Note) [PyTorch APIs for pooling layers](#)
- Its parameters and working is similar to a convolutional layer.
 - In PyTorch, `stride` is assigned as `kernel_size` if it is not given.




Pooling Layer

- Why is pooling important?
 - It reduces the network size. → Less time/space complexity
 - It provide larger **receptive fields** with a limited size of kernels.




Dropout


- **Dropout** is one of regularization methods against overfitting.
- It randomly drops a ratio of hidden units during training.
 - It prevents hidden units from co-adaptation.
 - e.g. When some hidden units has high weights, the others in the same layers have little contribution (low weight values, rarely training) to their output.



Standard



Dropout





DropConnect

Batch Normalization

- **Batch normalization (BatchNorm; BN)** is a technique for making neural network training more stable and smooth.
- It normalizes activations using the mean and variance of a mini-batch, and then learns how to scale and shift them. (N : batch size)


$$\text{BN}(x_{n,c,h,w}) = \gamma_c \cdot \left(\frac{x_{n,c,h,w} - \mu_c}{\sqrt{\sigma_c^2 + \epsilon}} \right) + \beta_c \quad \text{where} \quad \mu_c = \frac{1}{HWN} \sum_{n=1}^N \sum_{h=1}^H \sum_{w=1}^W x_{n,c,h,w} \quad \text{and} \quad \sigma_c^2 = \frac{1}{HWN} \sum_{i=1}^N \sum_{h=1}^H \sum_{w=1}^W (x_{n,c,h,w} - \mu_c)^2$$

- It often allows the network to use larger learning rates and converge faster.





Skip Connection


- Motivation) Why is a deeper network worse?



- A **skip connection** is a shortcut connection that bypasses several layers which contain nonlinearities (e.g. ReLU) and batch normalization.
 - Alias: *Residual connection* (used in *residual neural network*, *ResNet*)
 - It resolves the vanishing gradient problem.




gradient highway




Skip Connection

- Visualizing the loss landscape




(a) without skip connections



(b) with skip connections

Figure 1: The loss surfaces of ResNet-56 with/without skip connections. The proposed filter normalization scheme is used to enable comparisons of sharpness/flatness between the two figures.



Skip Connection

- Biological analogue [\[Wikipedia\]](#)
 - Cortical layer VI neurons @ [the cerebral cortex](#)





Image: [Wikipedia](#)

MNIST Dataset

- The [MNIST dataset](#) is a large database of handwritten digits. [\[Wikipedia\]](#)
- It was constructed by “re-mixing” the NIST’s original datasets.
 - Full name: *Modified* National Institute of Standards and Technology Database
- Specification
 - Classes: 10 (0, 1, 2, ..., 9)
 - Images: 28 x 28 (8-bit gray scale)
 - The number of data: 60,000 for training and 10,000 for test



Practice) Loading the MNIST Dataset

```
import torchvision
import matplotlib.pyplot as plt
```

Note) You can download the MNIST dataset through its mirror.

- Reference: <https://stackoverflow.com/questions/66577151/http-error-when-trying-to-download-mnist-data>

```
torchvision.datasets.MNIST.resources = [
    ('https://oss-ci-datasets.s3.amazonaws.com/mnist/train-images-idx3-ubyte.gz', 'f68b3c2dcbeaaa9fbdd348bbdeb94873'),
    ('https://oss-ci-datasets.s3.amazonaws.com/mnist/train-labels-idx1-ubyte.gz', 'd53e105ee54ea40749a09fcbcd1e9432'),
    ('https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-images-idx3-ubyte.gz', '9fb629c4189551a2d022fa330f9573f3'),
    ('https://oss-ci-datasets.s3.amazonaws.com/mnist/t10k-labels-idx1-ubyte.gz', 'ec29112dd5afa0611ce80d1b7f02629c')
]
```

Load the MNIST dataset


```
DATA_PATH = './data'
```

```
data_train = torchvision.datasets.MNIST(DATA_PATH, train=True, download=True)
```

```
data_valid = torchvision.datasets.MNIST(DATA_PATH, train=False)
```

Look inside of the dataset

```
print(data_train)          # ... 60000 ...
print(data_valid)          # ... 10000 ...
print(data_train.data.shape) # torch.Size([60000, 28, 28])
print(data_train.data.dtype) # torch.uint8
print(data_train.data[0,:,:]) # tensor([[0, 0, ..., ..., [166, 255, 247, ..., ...]])
plt.imshow(data_train.data[0,:,:), cmap='gray')
plt.show()
print(data_train.targets[0]) # Guess and check it!
```



Practice) Digit Classification with the MNIST Dataset (1/4)

```
# A four-layer CNN model
# - Try more or less layers, channels, and kernel size
# - Try to apply batch normalization (e.g. 'nn.BatchNorm' and 'nn.BatchNorm2d')
# - Try to apply skip connection (used in ResNet)
```

```
class MyCNN(nn.Module):
    def __init__(self):
        super(MyCNN, self).__init__()
        # Notation: (batch_size, channel, height, width)
        # Input: (batch_size, 1, 28, 28)
        # Layer1: conv (batch_size, 32, 28, 28)
        # pool (batch_size, 32, 14, 14)
        self.conv1 = nn.Conv2d(1, 32, kernel_size=3, stride=1, padding=1)
        self.pool1 = nn.MaxPool2d(kernel_size=2)

        # Layer2: conv (batch_size, 64, 14, 14)
        # pool (batch_size, 64, 7, 7)
        self.conv2 = nn.Conv2d(32, 64, kernel_size=3, stride=1, padding=1)
        self.pool2 = nn.MaxPool2d(kernel_size=2)
        self.drop2 = nn.Dropout(0.2)

        # Input: (batch_size, 64*7*7)
        # Layer3: fc (batch_size, 512)
        self.fc3 = nn.Linear(64*7*7, 512)
        self.drop3 = nn.Dropout(0.2)


        # Layer4: fc (batch_size, 10)
        self.fc4 = nn.Linear(512, 10)
```

```
def forward(self, x):
    x = F.relu(self.conv1(x))
    x = self.pool1(x)

    x = F.relu(self.conv2(x))
    x = self.pool2(x)
    x = self.drop2(x)
    x = torch.flatten(x, 1)

    x = F.relu(self.fc3(x))
    x = self.drop3(x)

    x = F.log_softmax(self.fc4(x), dim=1)
    return x
```



Practice) Digit Classification with the MNIST Dataset (2/4)

```
if __name__ == '__main__':
    # 0. Preparation
    torch.manual_seed(RANDOM_SEED)
    if USE_CUDA:
        torch.cuda.manual_seed_all(RANDOM_SEED)
    dev = torch.device('cuda' if USE_CUDA else 'cpu')

    # 1. Load the MNIST dataset
    preproc = torchvision.transforms.ToTensor()
    data_train = torchvision.datasets.MNIST(DATA_PATH, train=True, download=True, transform=preproc)
    data_valid = torchvision.datasets.MNIST(DATA_PATH, train=False, transform=preproc)
    loader_train = torch.utils.data.DataLoader(data_train, **DATA_LOADER_PARAM)
    loader_valid = torch.utils.data.DataLoader(data_valid, **DATA_LOADER_PARAM)

    # 2. Instantiate a model, loss function, and optimizer
    model = MyCNN().to(dev)
    loss_func = F.cross_entropy
    optimizer = torch.optim.SGD(model.parameters(), **OPTIMIZER_PARAM)
    scheduler = torch.optim.lr_scheduler.StepLR(optimizer, **SCHEDULER_PARAM)

    # 3.1. Train the model
    loss_list = []
    start = time.time()
    for epoch in range(1, EPOCH_MAX + 1):
        train_loss = train(model, loader_train, loss_func, optimizer)
        valid_loss, valid_accuracy = evaluate(model, loader_valid, loss_func)
        scheduler.step()

        loss_list.append([epoch, train_loss, valid_loss, valid_accuracy])
        if epoch % EPOCH_LOG == 0:
            elapse = (time.time() - start) / 60
            print(f'{epoch:>6} ({elapse:>6.2f} min), TrLoss={train_loss:.6f}, VaLoss={valid_loss:.6f}, VaAcc={valid_accuracy:.3f}, lr={scheduler.get_last_lr()}')
    elapse = (time.time() - start) / 60

    # 3.2. Save the trained model if necessary
    if SAVE_MODEL:
        torch.save(model.state_dict(), SAVE_MODEL)
```

Note) SCHEDULER_PARAM = { 'step_size': 10, 'gamma': 0.5 }

Practice) Digit Classification with the MNIST Dataset (3/4)

4.1. Visualize the loss curves

```
plt.title(f'Training and Validation Losses (time: {elapsed:.2f} [min] @ CUDA: {USE_CUDA})')
loss_array = np.array(loss_list)
plt.plot(loss_array[:,0], loss_array[:,1], label='Training Loss')
plt.plot(loss_array[:,0], loss_array[:,2], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss values')
plt.xlim(loss_array[0,0], loss_array[-1,0])
plt.grid()
plt.legend()
plt.show()
```





4.2. Visualize the confusion matrix

```
predicts = [predict(datum, model) for datum in data_valid.data]
conf_mat = metrics.confusion_matrix(data_valid.targets, predicts)
conf_fig = metrics.ConfusionMatrixDisplay(conf_mat)
conf_fig.plot()
```

5. Test your image

```
print(predict(data_train.data[0], model)) # 5
with PIL.Image.open('data/cnn_mnist_test.png').convert('L') as image:
    print(predict(image, model)) # 3
```

data_train.data[0] 'data/cnn_mnist_test.png'



Practice) Digit Classification with the MNIST Dataset (4/4)

Predict a digit using the given model

```
def predict(image, model):  
    model.eval()  
    with torch.no_grad():  
        # Convert the given image to its 1 x 1 x 28 x 28 tensor  
        if type(image) is torch.Tensor:  
            tensor = image.type(torch.float) / 255 # Normalize to [0, 1]  
        else:  
            tensor = 1 - TF.to_tensor(image) # Invert (white to black)  
        if tensor.ndim < 3:  
            tensor = tensor.unsqueeze(0)  
        if tensor.shape[0] == 3:  
            tensor = TF.rgb_to_grayscale(tensor) # Make grayscale  
        tensor = TF.resize(tensor, 28) # Resize to 28 x 28  
        dev = next(model.parameters()).device  
        tensor = tensor.unsqueeze(0).to(dev) # Add one more dim  
  
        output = model(tensor)  
        digit = torch.argmax(output, dim=1)  
        return digit.item()
```


Practice) Loading My Network and Testing My Image

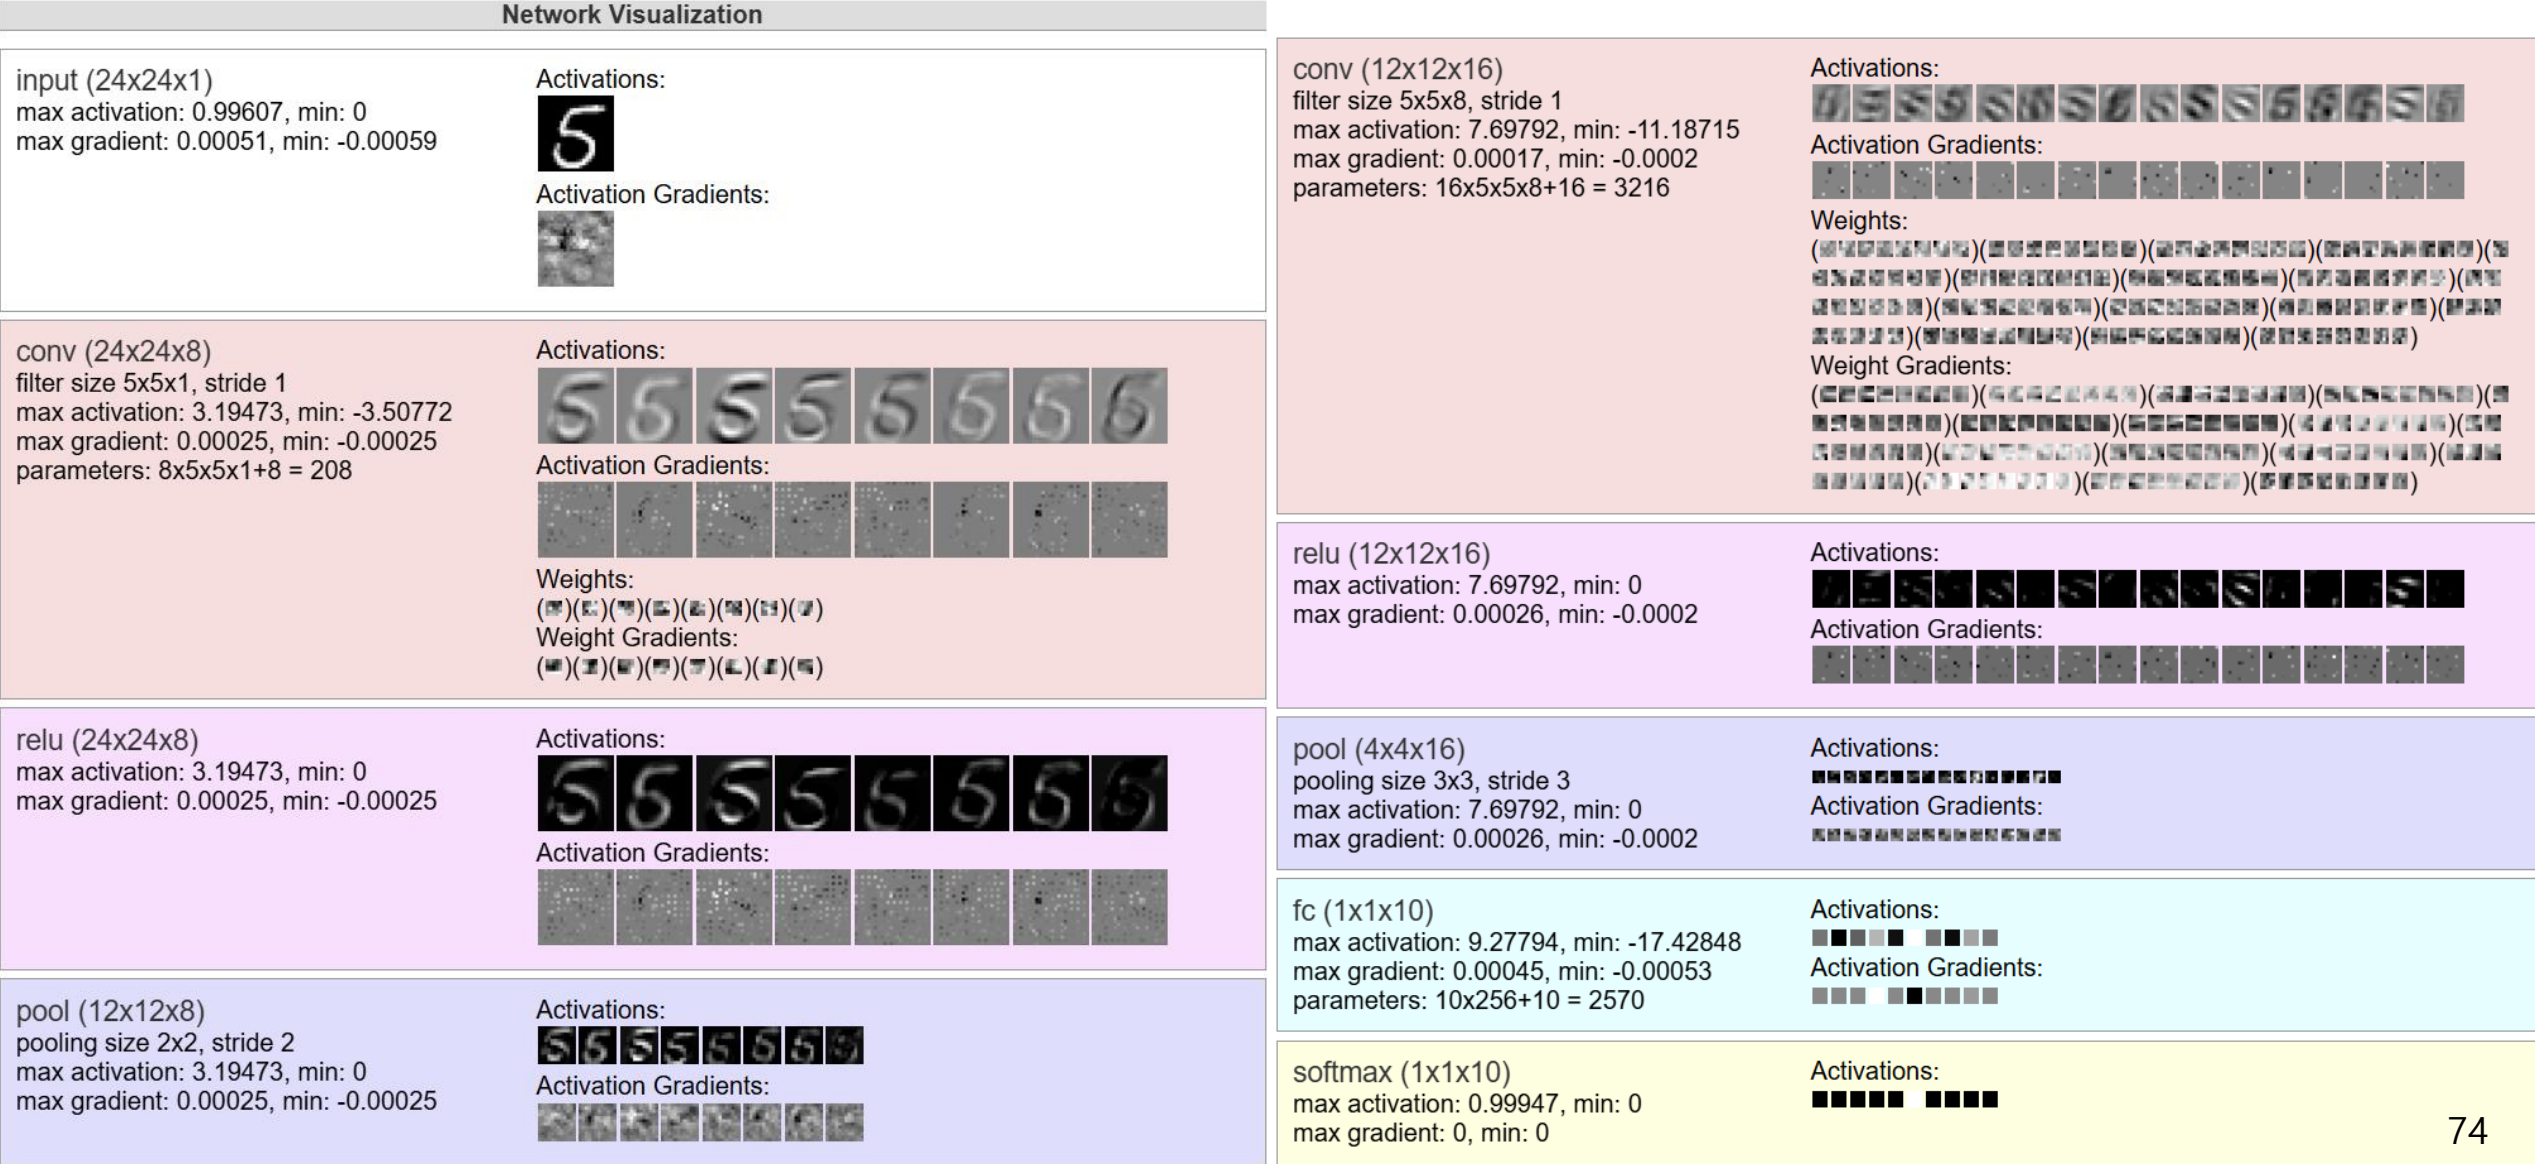
```
import torch, PIL
from cnn_mnist import MyCNN, predict

# Load a model
model = MyCNN()
model.load_state_dict(torch.load('cnn_mnist.pt'))



# Test the model
with PIL.Image.open('data/cnn_mnist_test.png').convert('L') as image:
    print(predict(image, model))
```

Practice) Digit Classification with the MNIST Dataset (Visualization)

- Example) ConvNetJS [MNIST demo](#)



Practice) Visualizing Learning Rate Schedulers (1/2)



```
import torch
import torch.optim.lr_scheduler as lr_scheduler
import matplotlib.pyplot as plt
```

```
base_lr = 1.
epoch_max = 100
schedulers = [
    {'name': 'StepLR',          'class': lr_scheduler.StepLR,
      'param': {'step_size': 10, 'gamma': 0.5}},
    {'name': 'LambdaLR',       'class': lr_scheduler.LambdaLR,
      'param': {'lr_lambda': lambda epoch: 0.5*(epoch / 10)}},
    {'name': 'MultiStepLR',     'class': lr_scheduler.MultiStepLR,
      'param': {'milestones': [40, 60, 70, 75, 80, 85, 90, 95], 'gamma': 0.5}},
    {'name': 'OneCycleLR',      'class': lr_scheduler.OneCycleLR,
      'param': {'max_lr': base_lr, 'total_steps': epoch_max}},
    {'name': 'CosAnnealing',    'class': lr_scheduler.CosineAnnealingWarmRestarts,
      'param': {'T_0': 50}},
    # Try more learning rate schedulers
]
```

Practice) Visualizing Learning Rate Schedulers (2/2)

```
for sch in schedulers:
    x = torch.tensor(1., requires_grad=True)           # A dummy parameter
    optimizer = torch.optim.SGD([x], lr=base_lr)        # Instantiate an optimizer
    scheduler = sch['class'](optimizer, **sch['param'])  # Instantiate a LR scheduler
    lr_values = []
    for i in range(epoch_max):
        optimizer.step()
        scheduler.step()
        lr_values.append(optimizer.param_groups[0]['lr'])
    plt.plot(range(epoch_max), lr_values, label=sch['name'])

plt.xlabel('Epochs')
plt.ylabel('Learning Rate')
plt.legend()
plt.show()
```

Practice) Different Styles for NN Classes (1/2)

My style

```
class MyCNN(nn.Module):
    def __init__(self):
        super(MyCNN, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1, 1)
        self.pool1 = nn.MaxPool2d(2)

        self.conv2 = nn.Conv2d(32, 64, 3, 1, 1)
        self.pool2 = nn.MaxPool2d(2)
        self.drop2 = nn.Dropout(0.2)

        self.fc3 = nn.Linear(64*7*7, 512)
        self.drop3 = nn.Dropout(0.2)

        self.fc4 = nn.Linear(512, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool1(x)

        x = F.relu(self.conv2(x))
        x = self.pool2(x)
        x = self.drop2(x)
        x = torch.flatten(x, 1)

        x = F.relu(self.fc3(x))
        x = self.drop3(x)

        x = F.log_softmax(self.fc4(x), dim=1)
        return x
```

Functional-oriented style

```
class MyCNN_Functional(nn.Module):
    def __init__(self):
        super(MyCNN_FStyle, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1, 1)
        self.fc1 = nn.Linear(64*7*7, 512)
        self.fc2 = nn.Linear(512, 10)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = F.max_pool2d(x, 2)

        x = F.relu(self.conv2(x))
        x = F.max_pool2d(x, 2)
        x = F.dropout(x, 0.2, self.training)
        x = torch.flatten(x, 1)

        x = F.relu(self.fc1(x))
        x = F.dropout(x, 0.2, self.training)

        x = F.log_softmax(self.fc2(x), dim=1)
        return x
```

Practice) Different Styles for NN Classes (2/2)

Object-oriented style

```
class MyCNN_Object(nn.Module):
    def __init__(self):
        super(MyCNN_ObjStyle, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1, 1)
        self.relu1 = nn.ReLU()
        self.pool1 = nn.MaxPool2d(2)

        self.conv2 = nn.Conv2d(32, 64, 3, 1, 1)
        self.relu2 = nn.ReLU()
        self.pool2 = nn.MaxPool2d(2)
        self.drop2 = nn.Dropout(0.2)

        self.fc3 = nn.Linear(64*7*7, 512)
        self.relu3 = nn.ReLU()
        self.drop3 = nn.Dropout(0.2)

        self.fc4 = nn.Linear(512, 10)
        self.smax4 = nn.LogSoftmax(dim=1)

    def forward(self, x):
        x = self.conv1(x)
        x = self.relu1(x)
        x = self.pool1(x)

        x = self.conv2(x)
        x = self.relu2(x)
        x = self.pool2(x)
        x = self.drop2(x)
        x = torch.flatten(x, 1)

        x = self.fc3(x)
```

Layer-oriented style

```
class MyCNN_Layer(nn.Module):
    def __init__(self):
        super(MyCNN_SeqStyle, self).__init__()
        self.layer1 = nn.Sequential(
            nn.Conv2d(1, 32, 3, 1, 1),
            nn.ReLU(),
            nn.MaxPool2d(2))

        self.layer2 = nn.Sequential(
            nn.Conv2d(32, 64, 3, 1, 1),
            nn.ReLU(),
            nn.MaxPool2d(2),
            nn.Dropout(0.2))

        self.layer3 = nn.Sequential(
            nn.Linear(64*7*7, 512),
            nn.ReLU(),
            nn.Dropout(0.2))


        self.layer4 = nn.Sequential(
            nn.Linear(512, 10),
            nn.LogSoftmax(dim=1))

    def forward(self, x):
        x = self.layer1(x)
        x = self.layer2(x)
        x = torch.flatten(x, 1)
        x = self.layer3(x)
        x = self.layer4(x)
        return x
```


Table of Contents

- Introduction
- PyTorch
- Neural Network (NN)
- Convolutional Neural Network (CNN)
- **Recurrent Neural Network (RNN)**
 - Recurrent neural network (RNN) unit
 - Long short-term memory (LSTM) unit
 - Gated recurrent unit (GRU)
 - Example) Name2Lang Classification with a Character-level RNN
- **Transformer**

Recurrent Neural Network (RNN)




- A **recurrent neural network** (shortly *RNN*) is a NN with a **loop**.
- A RNN can preserve a *memory* or (*hidden*) *state* or *information* inside.
 - Note) ~ sequential logic circuits (vs. combinatorial logic circuits) in digital circuits




Recurrent Neural Network (RNN)

- Multi-layer RNNs




Recurrent Neural Network (RNN)

- Memory in a RNN allows it to deal with *sequential* or *temporal* data.
- A **RNN** can deal with **various lengths of input vectors** and **output vectors** by attaching more hidden/output layers **at its end**.




Recurrent Neural Network (RNN)

- The **vanishing gradient problem** in training a long sequence data
 - A vanilla RNN cannot deal with long-term dependency.
 - e.g. Guessing the last word, "I grew up in France ... I speak fluent French."




Long Short-Term Memory (LSTM)



- A **long short-term memory** (shortly *LSTM*) is a recurrent neural network unit to deal with the vanishing gradient problem of a vanilla RNN. [\[Wikipedia\]](#)
- It is composed of a **forget gate**, an **input gate**, a **cell**, and an **output gate**.


forget gate	How much forget c_{t-1} to the cell state	$f_t = \sigma(W_f x_t + U_f h_{t-1} + b_f)$
input gate	How much write $\tilde{c}_t$ to the cell state	$i_t = \sigma(W_i x_t + U_i h_{t-1} + b_i)$
cell	The intermediate cell state	$\tilde{c}_t = \tanh(W_o x_t + U_o h_{t-1} + b_o)$
output gate	How much reveal c_t to the hidden state	$o_t = \sigma(W_o x_t + U_o h_{t-1} + b_o)$



Gated Recurrent Unit (GRU)

- A **gated recurrent unit** (shortly *GRU*) is a LSTM variant with a simplified structure, but it still has similar performance. [\[Wikipedia\]](#)

reset gate	How much forget h_{t-1} to the hidden state	$r_t = \sigma(W_r x_t + U_r h_{t-1} + b_r)$
update gate	How much write $\tilde{h}_t$ to the hidden state	$z_t = \sigma(W_z x_t + U_z h_{t-1} + b_z)$
	The intermediate hidden state	$\tilde{h}_t = \tanh(W_h x_t + U_h (r_t \otimes h_{t-1}) + b_h)$



Practice) Name2Lang Classification with a Character-level RNN (1/7)

- Motivation: A registration form on Internet

First Name

Last Name

E-mail address

Country

- **Input: Name** (string; a sequence of characters)
 - e.g. 'Choi', 'Jane', 'Daniel', 'Chow', 'Tanaka', ...
- **Classes: 18 languages** (integer; 0-17)
 - e.g. Arabic, Chinese, Czech, Dutch, English, French, German, Greek, Iris, Italian, Japanese, Korean, Polish, Portuguese, Russian, Scottish, Spanish, Vietnamese
- The dataset and bottom-up implementation are available on [the PyTorch official tutorial](#).

Practice) Name2Lang Classification with a Character-level RNN (2/7)

- **Input: Name** (string; a sequence of characters)
 - e.g. 'Choi', 'Jane', 'Daniel', 'Chow', 'Tanaka', ...
- How to represent the **name**?
 - The name is represented using 57 characters (52 alphabets and 5 special letters).
 - e.g. 'a' (0x61), 'b' (0x62), 'c' (0x63), ..., '-' (0x45) in ASCII code
 - Each character is encoded as a 57-bit one-hot vector.

• e.g. a:

1	0	0	0	...
---	---	---	---	-----

b:

0	1	0	0	...
---	---	---	---	-----

c:

0	0	1	0	...
---	---	---	---	-----

Why one-hot encoding?

• e.g. 'Choi' represented as a 4 x 57 array


C:	0	...	0	0	0	...	0	0	...	1	0	...
h:	0	...	1	0	0	...	0	0	...	0	0	...
o:	0	...	0	0	0	...	1	0	...	0	0	...
i:	0	...	0	1	0	...	0	0	...	0	0	...

Index) 7 8 14 28

Practice) Name2Lang Classification with a Character-level RNN (3/7)

```
# A simple RNN model
# - Try a different RNN unit such LSTM and GRU
# - Try less or more hidden units
# - Try more layers (e.g. 'num_layers=2') and dropout (e.g. 'dropout=0.4')
class MyRNN(nn.Module):
    def init(self, input_size, output_size):
        super(MyRNN, self).init()
        self.rnn = torch.nn.RNN(input_size, 128)
        self.fc = torch.nn.Linear(128, output_size)

    def forward(self, x):
        output, hidden = self.rnn(x)
        x = self.fc(output[-1]) # Use the output of the last sequence
        return x
```



Practice) Name2Lang Classification with a Character-level RNN (4/7)

```
# Convert Unicode to ASCII
# e.g. Ślusàrski to Slusarski
def unicode2ascii(s):
    return ''.join(
        c for c in unicodedata.normalize('NFD', s)
        if unicodedata.category(c) != 'Mn' and c in LETTER_DICT)

# Read raw files which contain names belong to each language
# Note) Each filename is used as its target's name.
def load_name_dataset(files):
    data = []
    targets = []
    target_names = []
    for idx, filename in enumerate(files):
        lang = os.path.splitext(os.path.basename(filename))[0]
        names = open(filename, encoding='utf-8').read().strip().split('\n')

        data += [unicode2ascii(name) for name in names]
        targets += [idx] * len(names)
        target_names.append(lang)
    return data, targets, target_names

# Transform the given text to its one-hot encoded tensor
# Note) Tensor size: len(text) x 1 x len(LETTER_DICT)
#               sequence_length x batch_size x input_size
def text2onehot(text, device='cpu'):
    tensor = torch.zeros(len(text), 1, len(LETTER_DICT), device=device)
    for idx, letter in enumerate(text):
        tensor[idx][0][LETTER_DICT.find(letter)] = 1
    return tensor
```

 Arabic.txt	13KB
 Chinese.txt	2KB
 Czech.txt	4KB
 Dutch.txt	3KB
 English.txt	27KB
 French.txt	3KB
 German.txt	6KB
 Greek.txt	2KB
 Irish.txt	2KB
 Italian.txt	6KB
 Japanese.txt	8KB
 Korean.txt	1KB
 Polish.txt	2KB
 Portuguese.txt	1KB
 Russian.txt	84KB
 Scottish.txt	1KB
 Spanish.txt	3KB
 Vietnamese.txt	1KB

Practice) Name2Lang Classification with a Character-level RNN (5/7)

```
if __name__ == '__main__':
    # 0. Preparation
    # 1. Load the name2lang dataset
    data, targets, target_names = load_name_dataset(glob.glob(DATA_PATH))
    data_train = [(text2onehot(data[i], device=dev), torch.LongTensor(
        [targets[i]]).to(dev)) for i in range(len(data))]
    random.shuffle(data_train)

    # 2. Instantiate a model, loss function, and optimizer
    # 3.1. Train the model
    # 3.2. Save the trained model if necessary

    # 4.1. Visualize the loss curves
    # 4.2. Visualize the confusion matrix
    predicts = [predict(datum, model) for datum in data]
    conf_mat = sklearn.metrics.confusion_matrix(targets, predicts, normalize='true')
    plt.imshow(conf_mat)
    plt.xlabel('Predicted label')
    plt.ylabel('True label')
    plt.gca().set_xticklabels([''] + target_names, rotation=90)
    plt.gca().set_yticklabels([''] + target_names)
    plt.gca().xaxis.set_major_locator(ticker.MultipleLocator(1))
    plt.gca().yaxis.set_major_locator(ticker.MultipleLocator(1))
    plt.show()

    # 5. Test your texts
    report_predict('Choi', model, target_names)
    report_predict('Jane', model, target_names)
    report_predict('Daniel', model, target_names)
    report_predict('Chow', model, target_names)
    report_predict('Tanaka', model, target_names)
```

Practice) Name2Lang Classification with a Character-level RNN (6/7)


* Name: **Choi**
1. **Korean** : 50.5 %
2. Chinese : 36.1 %
3. Vietnamese: 11.8 %
4. Russian : 0.9 %
5. Arabic : 0.5 %

* Name: **Jane**
1. **English** : 62.4 %
2. German : 10.5 %
3. Korean : 9.0 %
4. Chinese : 7.0 %
5. Dutch : 6.2 %

* Name: **Daniel**
1. **English** : 48.4 %
2. French : 20.6 %
3. Czech : 13.1 %
4. Russian : 6.6 %
5. Portuguese: 3.1 %

* Name: **Chow**
1. **Korean** : 69.8 %
2. Chinese : 16.0 %
3. English : 4.7 %
4. Vietnamese: 4.3 %
5. Russian : 3.1 %

* Name: **Tanaka**
1. **Japanese** : 84.4 %
2. Russian : 14.0 %
3. Czech : 1.3 %
4. English : 0.1 %
5. Irish : 0.1 %



Practice) Name2Lang Classification with a Character-level RNN (7/7)

Predict the best result of the given text

```
def predict(text, model):
    model.eval()
    with torch.no_grad():
        dev = next(model.parameters()).device
        text_tensor = text2onehot(text, dev) # Convert text to one-hot vectors
        output = model(text_tensor)[0]      # Get the last output
        lang = torch.argmax(output)         # Get the best among 18 classes
    return lang.item()
```

Predict and report top-k results of the given text

```
def report_predict(text, model, target_names, n_predict=5):
    print(f'* Name: {text}')
    model.eval()
    with torch.no_grad():
        dev = next(model.parameters()).device
        text_tensor = text2onehot(text, dev)
        output = model(text_tensor)[0]
        prob = nn.functional.softmax(output, dim=0) # Make output as probability
        top_val, top_idx = prob.topk(n_predict)     # Get top-k among 18 classes
        for i in range(len(top_val)):
            print(f' {i+1}. {target_names[top_idx[i].item()]:<10}:
{top_val[i]*100:4.1f} %')
```




Table of Contents

- Introduction
- PyTorch
- Neural Network (NN)
- Convolutional Neural Network (CNN)
- Recurrent Neural Network (RNN)
- **Transformer**
 - Word Embedding
 - Transformer

Word Representations

- **One-hot encoding** represents each word as a binary one-hot vector in which exactly one element is 1 and all other elements are 0.
 - e.g. apple: [1, 0, 0, 0, ...], bear: [0 1, 0, 0, ...], cat: [0, 0, 1, 0, ...], dog: [0, 0, 0, 1, ...], ...
 - The vocabulary size is equal to the dimension size of the one-hot vector.
 - Limitations
 - **Sparse representation** → very high-dimensional feature space
 - Large memory and computational requirements
 - Potential risk of **overfitting**, especially when the training data are limited
 - **Equal distance** between any two distinct words
 - Difficulty in capturing meaningful relationships between words

Word Representations


- **Word embedding** represents a word as a low-dimensional, dense vector of real numbers.

$$\text{word} \rightarrow \mathbf{e}_{\text{word}} \in \mathbb{R}^d \quad \text{where } d \text{ can be } 100, 128, 256, \dots$$

- Note) **Embedding** is a learned continuous vector representation in which geometric relationships reflect useful linguistic or semantic properties.
- The key idea is to learn word representations from the **contexts** in which words occur.
 - Words appearing in similar contexts tend to have similar vector representations.
 - e.g. $|\text{dog} - \text{cat}| < |\text{dog} - \text{apple}|$
 - e.g. $\text{puppy} - \text{dog} + \text{cat} = \text{kitten}$
- Approaches
 - Static word embedding: [Word2Vec](#) (2013), [GloVe](#) (Global Vectors; 2014), [fastText](#) (2016), ...
 - **Contextualized** word embeddings: [ELMo](#) (bidirectional LSTM; 2018), [BERT](#) (bidirectional Transformer; 2018), ...
 - e.g. I went to the **bank** to deposit some money. / We sat on the **bank** of the river. → **different vectors**

Word Embedding

- Online demo) [TensorFlow Embedding Projector](#): Word retrieval and visualization



Word Embedding

- Online demo) [WebVectors Semantic Calculator](#)
 - e.g. puppy – dog + cat = kitten

Semantic Calculator

Calculate ratios, such as «find a word D related to the word C in the same way as the word A is related to the word B». An example is given in the placeholder: which word is in the same relation to the word «father» as «daughter» is to «mother»? The answer is «son». [More on this...](#)



Word frequency

☒ High ☒ Medium ☐ Low

English Wikipedia

1. kitten NOUN 0.63



2. pup NOUN 0.52

3. piglet NOUN 0.51

4. pet NOUN 0.45



5. rabbit NOUN 0.45



- Parts of speech are shown for each word, if present in the model.

Choose the model:

☐ British National Corpus ☐ Google News ☒ English Wikipedia ☐ English Gigaword ☐ Norsk Aviskorpus


Show only:

☒ Nouns ☐ Proper names ☐ Verbs ☐ Adjectives ☐ Adverbs ☐ All of them

Calculate!



Transformer

- Paper) Vaswani et al., "Attention Is All You Need", NIPS, 2017 [arXiv](#)
 - e.g. Attention is like highlighting important parts of a book.




Transformer

- Paper) Vaswani et al., "Attention Is All You Need", NIPS, 2017 [arXiv](#)
 - e.g. Machine translation



Transformer


- **Transformer** is a family of NN architectures based on the **multi-head attention** mechanism.
 - Originally proposed as an **encoder-decoder** architecture for sequence-to-sequence tasks such as machine translation.
 - The encoder and decoder can be used independently depending on the task.
 - e.g. Encoder-only architecture: **BERT** → text classification
 - e.g. Decoder-only architecture: **GPT** → text generation
 - During training, **tokens in a sequence can be processed in parallel**, rather than sequentially as in RNNs.
 - This high degree of **parallelization** enables significantly **faster training** and **better scalability** than RNNs and their variants.
 - It has been widely applied to many other domains including non-sequential tasks (e.g. image classification).



Transformer

▪ Multi-head attention

- Multi-head? Use h parallel attention heads to capture different types of relationships among tokens
- Input: **Q**uery (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)
- **Scaled dot-product attention**
 - $\text{attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$
 - $\text{head}_i = \text{attention}(QW_i^Q, KW_i^K, VW_i^V)$
 - $\text{multihead}(Q, K, V) = \text{concat}(\text{head}_1, \text{head}_2, \dots, \text{head}_h) W^O$



Transformer

▪ Multi-head attention

- Multi-head? Use h parallel attention heads to capture different types of relationships among tokens
- Input: **Q**uery (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)
 - **Q**uery represents what information a token is looking for
 - **K**ey represents what information a token can be matched by
 - **V**alue represents the actual information to be retrieved or aggregated
- Analogue) Web searching

Query

 what is transformer in deep learning?



Wikipedia

[https://en.wikipedia.org/wiki/Transformer_\(deep_learning\)](https://en.wikipedia.org/wiki/Transformer_(deep_learning))

Key

Transformer (deep learning)

Value

In deep learning, the transformer is a family of artificial neural network architectures based on the multi-head attention mechanism, in which input data ...

History

Training

Architecture



GeeksforGeeks

<https://www.geeksforgeeks.org/machine-learning/geeksforgeeks-transformers-in-machine-learning/>

Key

Transformers in Machine Learning



Value

2026. 5. 11. — Transformer is a neural network architecture used for various machine learning tasks, especially in natural language processing and computer ...

Transformer

Multi-head attention



- Multi-head? Use h parallel attention heads to capture different types of relationships among tokens
- Input: **Q**uery (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)
- **Scaled dot-product attention**
 - $\text{attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$
 - $\text{head}_i = \text{attention}(QW_i^Q, KW_i^K, VW_i^V)$
 - $\text{multihead}(Q, K, V) = \text{concat}(\text{head}_1, \text{head}_2, \dots, \text{head}_h) W^O$



Transformer

Multi-head attention

- Multi-head? Use h parallel attention heads to capture different types of relationships among tokens
- Input: **Q**uery (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)
- **Scaled dot-product attention**
 - $\text{attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$
 - $\text{head}_i = \text{attention}(QW_i^Q, KW_i^K, VW_i^V)$
 - $\text{multihead}(Q, K, V) = \text{concat}(\text{head}_1, \text{head}_2, \dots, \text{head}_h) W^O$
 - Q) Why not compute attention directly?
(w/o queries and keys)





Transformer

Multi-head attention

- Multi-head? Use h parallel attention heads to capture different types of relationships among tokens
- Input: **Q**uery (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)

Scaled dot-product attention

- $\text{attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$ Note) $d_k = 2$
- $\text{head}_i = \text{attention}(QW_i^Q, KW_i^K, VW_i^V)$
- $\text{multihead}(Q, K, V) = \text{concat}(\text{head}_1, \text{head}_2, \dots, \text{head}_h) W^O$





Transformer

Multi-head attention

- Multi-head? Use h parallel attention heads to capture different types of relationships among tokens
- Input: **Q**uery (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)
- **Scaled dot-product attention**



- $\text{attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$ Note) $d_k = 2$
- $\text{head}_i = \text{attention}(QW_i^Q, KW_i^K, VW_i^V)$
- $\text{multihead}(Q, K, V) = \text{concat}(\text{head}_1, \text{head}_2, \dots, \text{head}_h) W^O$



Transformer

Multi-head attention

- Multi-head? Use h parallel attention heads to capture different types of relationships among tokens
- Input: **Q**uery (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)
- **Scaled dot-product attention**
 - $\text{attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$ Note) $d_k = 2$
 - $\text{head}_i = \text{attention}(QW_i^Q, KW_i^K, VW_i^V)$
 - $\text{multihead}(Q, K, V) = \text{concat}(\text{head}_1, \text{head}_2, \dots, \text{head}_h) W^O$





Transformer

Multi-head attention

- Multi-head? Use h parallel attention heads to capture different types of relationships among tokens
- Input: **Q**uery (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)
- **Scaled dot-product attention**


- $\text{attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) V$ Note) $d_k = 2$
- $\text{head}_i = \text{attention}(QW_i^Q, KW_i^K, VW_i^V)$
- $\text{multihead}(Q, K, V) = \text{concat}(\text{head}_1, \text{head}_2, \dots, \text{head}_h) W^O$



Transformer

▪ Positional encoding



- Q) How can the model capture the position of each word (token)?
 - e.g. I am a student vs. am I a student
- A) Add positional information to the token embeddings!



Transformer

▪ Positional encoding


- Q) How can the model capture the position of each word (token)?
 - e.g. I am a student vs. am I a student
- A) Add positional information to the token embeddings!



Transformer


▪ Positional encoding

- Q) How can the model capture the position of each word (token)?
 - e.g. I am a student vs. am I a student
- A) Add positional information to the token embeddings!
- Q) How are the positional encoding values determined?
 - $PE(pos, 2i) = \sin\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$
 - $PE(pos, 2i + 1) = \cos\left(\frac{pos}{10000^{\frac{2i}{d_{model}}}}\right)$
 - Example) d_{model} : 128, sequence length: 50 ([code available](#))
 - Note) Modern Transformer models often use learned positional embeddings, relative positional representations, or RoPE (Rotary Position Embedding) instead of fixed sinusoidal positional encoding..




Transformer

- **Transformer** is a family of NN architectures based on the **multi-head attention** mechanism.
 - Originally proposed as an **encoder-decoder** architecture for sequence-to-sequence tasks such as machine translation.



Transformer



- **Transformer** is a family of NN architectures based on the **multi-head attention** mechanism.
 - Originally proposed as an **encoder-decoder** architecture for sequence-to-sequence tasks such as machine translation.



Transformer

- **Transformer** is a family of NN architectures based on the **multi-head attention** mechanism.
 - Originally proposed as an **encoder-decoder** architecture for sequence-to-sequence tasks such as machine translation.

The output of the final encoder is provided as input to each decoder layer.




Transformer

Multi-head attention

- Input: **Q**uey (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)
- Types
 - **Encoder self-attention**: Q, K, V from encoders
 - **Masked decoder self-attention**: Q, K, V from decoders
 - **Encoder-decoder attention**: Q from decoders
K, V from the last encoder


Encoder Self-Attention:



Masked Decoder Self-Attention:



Encoder-Decoder Attention:




Transformer

▪ Multi-head attention

- Input: **Q**uery (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)
- Types
 - **Encoder self-attention**: Q, K, V from encoders
 - **Masked decoder self-attention**: Q, K, V from decoders
 - **Encoder-decoder attention**: Q from decoders
K, V from the last encoder

Encoder Self-Attention:




Masked Decoder Self-Attention:



Encoder-Decoder Attention:



Self-attention?





Transformer


Multi-head attention

- Input: **Q**uery (질의 in Korean), **K**ey (요약 in Korean), **V**alue (내용 in Korean)
- Types
 - Encoder self-attention**: Q, K, V from encoders
 - Masked decoder self-attention**: Q, K, V from decoders
 - Encoder-decoder attention**: Q from decoders
K, V from the last encoder

Note) **Mask padding tokens** to prevent them from affecting the model's computation




Encoder Self-Attention:



Masked Decoder Self-Attention:



Encoder-Decoder Attention:



Example) Transformer from Scratch (by [graykode](#))

```
if __name__ == '__main__':
    sentences = ['ich mochte ein bier P', 'S i want a beer', 'i want a beer E']

    # Transformer Parameters
    # Padding Should be Zero
    src_vocab = {'P': 0, 'ich': 1, 'mochte': 2, 'ein': 3, 'bier': 4}
    src_vocab_size = len(src_vocab)

    tgt_vocab = {'P': 0, 'i': 1, 'want': 2, 'a': 3, 'beer': 4, 'S': 5, 'E': 6}
    number_dict = {i: w for i, w in enumerate(tgt_vocab)}
    tgt_vocab_size = len(tgt_vocab)

    src_len = 5      # length of source
    tgt_len = 5      # length of target

    d_model = 512    # Embedding Size
    d_ff = 2048       # FeedForward dimension
    d_k = d_v = 64    # dimension of K(=Q), V
    n_layers = 6      # number of Encoder of Decoder Layer
    n_heads = 8       # number of heads in Multi-Head Attention

    model = Transformer()

    criterion = nn.CrossEntropyLoss()
    optimizer = optim.Adam(model.parameters(), lr=0.001)

    enc_inputs, dec_inputs, target_batch = make_batch(sentences)

    for epoch in range(20):
        optimizer.zero_grad()
        outputs, enc_self_attns, dec_self_attns, dec_enc_attns = model(enc_inputs, dec_inputs)
```

Example) Transformer from Scratch (by [graycode](#))



```
def make_batch(sentences):
    input_batch = [[src_vocab[n] for n in sentences[0].split()]]
    output_batch = [[tgt_vocab[n] for n in sentences[1].split()]]
    target_batch = [[tgt_vocab[n] for n in sentences[2].split()]]
    return torch.LongTensor(input_batch), torch.LongTensor(output_batch), torch.LongTensor(target_batch)
```

```
def get_sinusoid_encoding_table(n_position, d_model):
    def cal_angle(position, hid_idx):
        return position / np.power(10000, 2 * (hid_idx // 2) / d_model)
    def get_posi_angle_vec(position):
        return [cal_angle(position, hid_j) for hid_j in range(d_model)]

    sinusoid_table = np.array([get_posi_angle_vec(pos_i) for pos_i in range(n_position)])
    sinusoid_table[:, 0::2] = np.sin(sinusoid_table[:, 0::2]) # dim 2i
    sinusoid_table[:, 1::2] = np.cos(sinusoid_table[:, 1::2]) # dim 2i+1
    return torch.FloatTensor(sinusoid_table)
```

```
def get_attn_pad_mask(seq_q, seq_k):
    batch_size, len_q = seq_q.size()
    batch_size, len_k = seq_k.size()
    # eq(zero) marks PAD token as 1.
    pad_attn_mask = seq_k.data.eq(0).unsqueeze(1) # batch_size x 1 x len_k(=len_q), one is masking
    return pad_attn_mask.expand(batch_size, len_q, len_k) # batch_size x len_q x len_k
```

```
def get_attn_subsequent_mask(seq):
    attn_shape = [seq.size(0), seq.size(1), seq.size(1)]
    subsequent_mask = np.triu(np.ones(attn_shape), k=1)
    subsequent_mask = torch.from_numpy(subsequent_mask).byte()
    return subsequent_mask
```



Example) Transformer from Scratch (by [graycode](#))

```
class ScaledDotProductAttention(nn.Module):
    def __init__(self):
        super(ScaledDotProductAttention, self).__init__()

    def forward(self, Q, K, V, attn_mask):
        scores = torch.matmul(Q, K.transpose(-1, -2)) / np.sqrt(d_k) # [batch_size x n_heads x len_q(=len_k) x len_k(=len_q)]
        scores.masked_fill_(attn_mask, -1e9) # Fills elements of self tensor with value where mask is one.
        attn = nn.Softmax(dim=-1)(scores)
        context = torch.matmul(attn, V)
        return context, attn
```

Diagram illustrating the Scaled Dot Product Attention mechanism:

Query matrix Q (yellow, 4x2) is multiplied by the transpose of the Key matrix K^T (orange, 2x4). The result is passed through a softmax function, scaled by the square root of the key dimension $\sqrt{d_k}$. The resulting attention weights are then multiplied by the Value matrix V (green, 4x2) to produce the final Attention Value Matrix α (blue, 4x2).

$$\text{softmax} \left(\frac{Q \times K^T}{\sqrt{d_k}} \right) \times V = \text{Attention Value Matrix } \alpha$$


Example) Transformer from Scratch (by [graycode](#))

```
class MultiHeadAttention(nn.Module):
    def __init__(self):
        super(MultiHeadAttention, self).__init__()
        self.W_Q = nn.Linear(d_model, d_k * n_heads)
        self.W_K = nn.Linear(d_model, d_k * n_heads)
        self.W_V = nn.Linear(d_model, d_v * n_heads)
        self.linear = nn.Linear(n_heads * d_v, d_model)
        self.layer_norm = nn.LayerNorm(d_model)

    def forward(self, Q, K, V, attn_mask):
        # q: [batch_size x len_q x d_model], k: [batch_size x len_k x d_model], v: [batch_size x len_k x d_model]
        residual, batch_size = Q, Q.size(0)
        # (B, S, D) -proj-> (B, S, D) -split-> (B, S, H, W) -trans-> (B, H, S, W)
        q_s = self.W_Q(Q).view(batch_size, -1, n_heads, d_k).transpose(1,2) # q_s: [batch_size x n_heads x len_q x d_k]
        k_s = self.W_K(K).view(batch_size, -1, n_heads, d_k).transpose(1,2) # k_s: [batch_size x n_heads x len_k x d_k]
        v_s = self.W_V(V).view(batch_size, -1, n_heads, d_v).transpose(1,2) # v_s: [batch_size x n_heads x len_k x d_v]

        attn_mask = attn_mask.unsqueeze(1).repeat(1, n_heads, 1, 1) # attn_mask : [batch_size x n_heads x len_q x len_k]

        # context: [batch_size x n_heads x len_q x d_v], attn: [batch_size x n_heads x len_q(=len_k) x len_k(=len_q)]
        context, attn = ScaledDotProductAttention()(q_s, k_s, v_s, attn_mask)
        context = context.transpose(1, 2).contiguous().view(batch_size, -1, n_heads * d_v) # context: [batch_size x len_q x n_heads * d_v]
        output = self.linear(context)
        return self.layer_norm(output + residual), attn # output: [batch_size x len_q x d_model]
```




Example) Transformer from Scratch (by [graycode](#))

```
class EncoderLayer(nn.Module):
    def __init__(self):
        super(EncoderLayer, self).__init__()
        self.enc_self_attn = MultiHeadAttention()
        self.pos_ffn = PoswiseFeedForwardNet()

    def forward(self, enc_inputs, enc_self_attn_mask):
        enc_outputs, attn = self.enc_self_attn(enc_inputs, enc_inputs, enc_inputs, enc_self_attn_mask) # enc_inputs to same Q,K,V
        enc_outputs = self.pos_ffn(enc_outputs) # enc_outputs: [batch_size x len_q x d_model]
        return enc_outputs, attn

class Encoder(nn.Module):
    def __init__(self):
        super(Encoder, self).__init__()
        self.src_emb = nn.Embedding(src_vocab_size, d_model)
        self.pos_emb = nn.Embedding.from_pretrained(get_sinusoid_encoding_table(src_len+1, d_model), freeze=True)
        self.layers = nn.ModuleList([EncoderLayer() for _ in range(n_layers)])

    def forward(self, enc_inputs): # enc_inputs : [batch_size x source_len]
        enc_outputs = self.src_emb(enc_inputs) + self.pos_emb(torch.LongTensor([[1, 2, 3, 4, 0]]))
        enc_self_attn_mask = get_attn_pad_mask(enc_inputs, enc_inputs)
        enc_self_attns = []
        for layer in self.layers:
            enc_outputs, enc_self_attn = layer(enc_outputs, enc_self_attn_mask)
            enc_self_attns.append(enc_self_attn)
        return enc_outputs, enc_self_attns
```



Example) Transformer from Scratch (by [graycode](#))

```
class DecoderLayer(nn.Module):
    def __init__(self):
        super(DecoderLayer, self).__init__()
        self.dec_self_attn = MultiHeadAttention()
        self.dec_enc_attn = MultiHeadAttention()
        self.pos_ffn = PoswiseFeedForwardNet()


    def forward(self, dec_inputs, enc_outputs, dec_self_attn_mask, dec_enc_attn_mask):
        dec_outputs, dec_self_attn = self.dec_self_attn(dec_inputs, dec_inputs, dec_inputs, dec_self_attn_mask)
        dec_outputs, dec_enc_attn = self.dec_enc_attn(dec_outputs, enc_outputs, enc_outputs, dec_enc_attn_mask)
        dec_outputs = self.pos_ffn(dec_outputs)
        return dec_outputs, dec_self_attn, dec_enc_attn

class Decoder(nn.Module):
    def __init__(self):
        super(Decoder, self).__init__()
        self.tgt_emb = nn.Embedding(tgt_vocab_size, d_model)
        self.pos_emb = nn.Embedding.from_pretrained(get_sinusoid_encoding_table(tgt_len+1, d_model), freeze=True)
        self.layers = nn.ModuleList([DecoderLayer() for _ in range(n_layers)])

    def forward(self, dec_inputs, enc_inputs, enc_outputs): # dec_inputs : [batch_size x target_len]
        dec_outputs = self.tgt_emb(dec_inputs) + self.pos_emb(torch.LongTensor([[5, 1, 2, 3, 4]]))
        dec_self_attn_pad_mask = get_attn_pad_mask(dec_inputs, dec_inputs)
        dec_self_attn_subsequent_mask = get_attn_subsequent_mask(dec_inputs)
        dec_self_attn_mask = torch.gt((dec_self_attn_pad_mask + dec_self_attn_subsequent_mask), 0)

        dec_enc_attn_mask = get_attn_pad_mask(dec_inputs, enc_inputs)


        dec_self_attns, dec_enc_attns = [], []
        for layer in self.layers:
            dec_outputs, dec_self_attn, dec_enc_attn = layer(dec_outputs, enc_outputs, dec_self_attn_mask, dec_enc_attn_mask)
            dec_self_attns.append(dec_self_attn)
            dec_enc_attns.append(dec_enc_attn)
        return dec_outputs, dec_self_attns, dec_enc_attns
```



Example) Transformer from Scratch (by [graycode](#))

```
class Transformer(nn.Module):
    def __init__(self):
        super(Transformer, self).__init__()
        self.encoder = Encoder()
        self.decoder = Decoder()
        self.projection = nn.Linear(d_model, tgt_vocab_size, bias=False)

    def forward(self, enc_inputs, dec_inputs):
        enc_outputs, enc_self_attns = self.encoder(enc_inputs)
        dec_outputs, dec_self_attns, dec_enc_attns = self.decoder(dec_inputs, enc_inputs, enc_outputs)
        dec_logits = self.projection(dec_outputs) # dec_logits : [batch_size x src_vocab_size x tgt_vocab_size]
        return dec_logits.view(-1, dec_logits.size(-1)), enc_self_attns, dec_self_attns, dec_enc_attns
```



Example) Transformer from Scratch (by [graycode](#))

```
if __name__ == '__main__':
    sentences = ['ich mochte ein bier P', 'S i want a beer', 'i want a beer E']

    ...


    for epoch in range(20):
        optimizer.zero_grad()
        outputs, enc_self_attns, dec_self_attns, dec_enc_attns = model(enc_inputs, dec_inputs)
        loss = criterion(outputs, target_batch.contiguous().view(-1))
        print('Epoch:', '%04d' % (epoch + 1), 'cost =', '{:.6f}'.format(loss))
        loss.backward()
        optimizer.step()

    # Test
    predict, _, _, _ = model(enc_inputs, dec_inputs)
    predict = predict.data.max(1, keepdim=True)[1]
    print(sentences[0], '->', [number_dict[n.item()] for n in predict.squeeze()])

    print('first head of last state enc_self_attns')
    showgraph(enc_self_attns)

    print('first head of last state dec_self_attns')
    showgraph(dec_self_attns)

    print('first head of last state dec_enc_attns')
    showgraph(dec_enc_attns)
```



Summary

- **DNN:** Deep neural network
- **CNN:** Feedforward with convolution (and pooling)
- **RNN:** NN with a loop → state/memory → sequential/temporal data
- **Transformer:** Multi-head attention mechanism, positional encoding → sequential/temporal data

- **Issue #1) Vanishing gradient problem**
 - Activation functions such as ReLU
 - Skip connection ~ LSTM and GRU
- **Issue #2) Overfitting problem**
 - More data by data collection or data augmentation or data synthesis
 - Data separation (train/validation/test data), cross-validation, and early stopping
 - More simplified models (e.g. CNN ← weight sharing and local connectivity; a.k.a. inductive bias)
 - Loss functions with regularization terms
- Other improvement
 - Dropout
 - Batch normalization
 - ...

Further Information

- Natural Language Processing (NLP)
 - **Seq2seq** (2014), **attention mechanism** (2015)
 - **Transformer** (2017), **BERT** (Bidirectional Encoder Representations from Transformers), **GPT** (Generative Pretrained Transformer), ...
- Computer Vision
 - CNN backbone networks: **AlexNet** (2012), **VGGNet** (2014), **Inception Net** (2014), **ResNet** (2015), ...
 - CNN object detection networks
 - Two-stage detectors: **R-CNN** (2014), **Fast R-CNN** (2015), **Faster R-CNN** (2015), **Mask R-CNN** (2017), ...
 - One-stage detectors: **YOLO** (You Only Look Once; 2015), **SSD** (Single Shot MultiBox Detector; 2015), ...
 - Vision Transformers (ViT) and Vision Foundation Models (VFM)
 - **ViT** (Vision Transformer; 2020), **Swin Transformer** (2021), **DINO** (2021), ..., **SAM** (Segment Anything; 2023), ...
 - Generative models
 - **Autoencoder** (2006), ..., **GAN** (Generative Adversarial Networks; 2014), **CycleGAN** (2017), ...
 - **Diffusion model** (2020), **Stable Diffusion** (2022), ...
 - Others
 - **NeRF** (Neural Radiance Field; 2020), **CLIP** (Contrastive Language-Image Pre-Training; 2021), ...
- Others
 - Continual/incremental learning, ..., transfer learning/domain adaptation, ..., contrastive learning, ...
 - Network compression/pruning, knowledge distillation (2015), ..., ML model deployment, ...

How To Understand New Methods (or Read Papers)

- **Model**

- Network design
- Input and output
 - e.g. Normalization, positional encoding

- **Training**

- Dataset
 - e.g. Data augmentation, data synthesis
 - e.g. Data labeling, unsupervised or weakly-supervised approaches
- Objective functions
 - e.g. Regularization terms, [Huber loss](#) (for robust regression), [focal loss](#)
- Optimization
 - e.g. Optimization algorithms, hyperparameters